

Advanced Machine Learning

Imbalanced Learning via Cost-Sensitive Learning

Learning goals

- Learn the modelling approach via the cost matrix in cost-sensitive learning settings
- Understand the connection between cost-sensitive and imbalanced learning
- Get to know MetaCost as a general approach to make classifiers cost-sensitive

		Confusion matrix	
		True class	
		$y = 1$	$y = -1$
Pred. $\hat{y} = 1$ class $\hat{y} = -1$	$y = 1$	TP	FP
	$y = -1$	FN	TN

		Cost matrix	
		True class	
		$y = 1$	$y = -1$
Pred. $\hat{y} = 1$ class $\hat{y} = -1$	$y = 1$	$C(1, 1)$	$C(1, -1)$
	$y = -1$	$C(-1, 1)$	$C(-1, -1)$

COST-SENSITIVE LEARNING: IN A NUTSHELL

- Cost-sensitive learning is a learning paradigm, where different (mis-)classification costs are taken into consideration and the learner seeks to minimize the total costs in expectation.
- Thus, the major difference to the “classical” cost-insensitive learning setting is that cost-sensitive learning deals differently with misclassifications. Most of the algorithms of the latter kind assume that the data sets are balanced, and all errors have the same cost.
- This is motivated by a plethora of real-world applications, where different costs between misclassifications are present:
 - Medicine — Misdiagnosing a cancer patient as healthy vs. misdiagnosing a healthy patient as having cancer (and then check again).
 - Tracking criminals — Classify an innocent person as a terrorist vs. overlooking a terrorist.
 - (Extreme) Weather prediction — Incorrectly predicting that no hurricane occurs vs. predicting a strong wind as a hurricane.
 - Credit granting — Lending to a risky client vs. not lending to a trustworthy client.
 - . . .
- In all these examples the costs of a false negative is much higher than the costs of a false positive.

COST MATRIX

- In cost-sensitive learning we are provided with a cost matrix **C** of the form

		True Class y			
		1	2	...	g
Classification \hat{y}	1	$C(1, 1)$ (True 1's)	$C(1, 2)$ (False 1's for 2's)	...	$C(1, g)$ (False 1's for g 's)
	2	$C(2, 1)$ (False 2's for 1's)	$C(2, 2)$ (True 2's)	...	$C(2, g)$ (False 2's for g 's)
	\vdots	\vdots	\vdots	...	\vdots
	\vdots	\vdots	\vdots	...	\vdots
	g	$C(g, 1)$ (False g 's for 1's)	$C(g, 2)$ (False g 's for 2's)	...	$C(g, g)$ (True g 's)

- Here, $C(i, j)$ is the cost of classifying j as i , which in the cost-insensitive learning case is simply $C(i, j) = \mathbb{1}_{[i \neq j]}$, i.e., each misclassification has the same cost of 1.
- Sometimes the cost matrix is provided by the application at hand, e.g. in the credit granting example, or provided by a domain expert. In many cases, however, the cost matrix needs to be estimated. This is usually done by using a heuristic or by learning a proper cost matrix from the training data.
- The cost matrix is essential for the learning process, as
 - 1 too low costs might not change the decision boundaries significantly leading to (still) costly predictions,
 - 2 too high costs might harm the generalization capability of the classifier on costly classes.

COST MATRIX FOR IMBALANCED LEARNING

- For imbalanced data sets biases towards majority classes can be enlarged or even too strong biases towards the minority can be created if the cost matrix is set inappropriately.
- A common heuristic for imbalanced data sets is to use
 - the imbalance ratio between majority and minority classes for misclassifying a minority class j as a majority class i , i.e., $C(i, j) = \frac{n_i}{n_j}$ for classes i and j such that $n_j \ll n_i$,
 - and costs of 1 for misclassifying a majority class j as a minority class i , i.e., $C(i, j) = 1$ for classes i and j such that $n_i \ll n_j$.
 - Usually, the costs of a correct classification is set to 0, which is justified also from a theoretical point of view as we will see later.

- In an imbalanced binary classification problem we obtain the following cost matrix using this heuristic:

	True class	
	$y = 1$	$y = -1$
Pred. $\hat{y} = 1$	0	1
class $\hat{y} = -1$	$\frac{n_-}{n_+}$	0

- Thus, this heuristic is consistent with the general real case that the cost of false negatives is much higher than the cost of false positives.

MINIMUM EXPECTED COST PRINCIPLE

- Suppose we are provided with a cost matrix \mathbf{C} and also we have knowledge of the true posterior distribution $p(\cdot \mid \mathbf{x})$. The most natural way to classify a given feature \mathbf{x} is then by following the *minimum expected cost principle*.
- Minimum expected cost principle: Use the class for prediction with the smallest expected costs, where the expected costs of a class $i \in \{1, \dots, g\}$ is

$$\mathbb{E}_{J \sim p(\cdot \mid \mathbf{x})}(C(i, J)) = \sum_{j=1}^g p(j \mid \mathbf{x}) C(i, j).$$

- Thus, if we have a classifier f which uses a probabilistic score function $\pi : \mathcal{X} \rightarrow [0, 1]^g$ with $\pi(\mathbf{x}) = (\pi(\mathbf{x})_1, \dots, \pi(\mathbf{x})_g)^\top$ and $\sum_{j=1}^g \pi(\mathbf{x})_j = 1$ for the classification, then one can easily modify f to take the expected costs into account:

$$\tilde{f}(\mathbf{x}) = \arg \min_{i=1, \dots, g} \sum_{j=1}^g \pi(\mathbf{x})_j C(i, j).$$

MINIMUM EXPECTED COST PRINCIPLE: BINARY CASE

- In the binary classification setting (i.e., $\mathcal{Y} = \{-1, 1\}$) the minimum expected costs principle translates to predict the positive class if

$$\begin{aligned}\mathbb{E}_{J \sim p(\cdot | \mathbf{x})}(C(1, J)) &\leq \mathbb{E}_{J \sim p(\cdot | \mathbf{x})}(C(-1, J)) \\ \Leftrightarrow p(-1 | \mathbf{x})C(1, -1) + p(1 | \mathbf{x})C(1, 1) \\ &\leq p(-1 | \mathbf{x})C(-1, -1) + p(1 | \mathbf{x})C(-1, 1) \\ \Leftrightarrow p(-1 | \mathbf{x})(C(1, -1) - C(-1, -1)) &\leq p(1 | \mathbf{x})(C(-1, 1) - C(1, 1))\end{aligned}$$

- Note that the decision for predicting the positive class does not change if we use instead of \mathbf{C} the simpler cost matrix \mathbf{C}_{simple} , where

- $C_{simple}(-1, -1) = C_{simple}(1, 1) = 0$
- $C_{simple}(1, -1) = C(1, -1) - C(-1, -1)$
- $C_{simple}(-1, 1) = C(-1, 1) - C(1, 1).$

MINIMUM EXPECTED COST PRINCIPLE: BINARY CASE

- Thus, one can assume without loss of generality that the cost matrix \mathbf{C}

	True class	
	$y = 1$	$y = -1$
Pred. $\hat{y} = 1$	$C(1, 1)$	$C(1, -1)$
class $\hat{y} = -1$	$C(-1, 1)$	$C(-1, -1)$

is of a simpler form \mathbf{C}_{simple} :

	True class	
	$y = 1$	$y = -1$
Pred. $\hat{y} = 1$	0	$C(1, -1) - C(-1, -1)$
class $\hat{y} = -1$	$C(-1, 1) - C(1, 1)$	0

- An analogous result can be shown for the multiclass setting.

MINIMUM EXPECTED COST PRINCIPLE: BINARY CASE

- With this simpler cost matrix (relabeling \mathbf{C} by \mathbf{C}_{simple}), the decision to predict the positive class boils down to

$$\begin{aligned} p(-1 | \mathbf{x})C(1, -1) &\leq p(1 | \mathbf{x})C(-1, 1) \\ \Leftrightarrow (1 - p(1 | \mathbf{x}))C(1, -1) &\leq p(1 | \mathbf{x})C(-1, 1) \\ \Leftrightarrow \underbrace{\frac{C(1, -1)}{C(1, -1) + C(-1, 1)}}_{=: c^*} &\leq p(1 | \mathbf{x}) \end{aligned}$$

- This yields the optimal threshold value c^* for probabilistic score classifiers, so that any probabilistic classifier f using a probabilistic score $\pi : \mathcal{X} \rightarrow [0, 1]$ can be modified to

$$\tilde{f}(\mathbf{x}) = 2 \cdot \mathbb{1}_{[\pi(\mathbf{x}) \geq c^*]} - 1.$$

METACOST: OVERVIEW

- Substantial work has gone into making individual algorithms cost-sensitive. A better solution would be to have a procedure that can convert a broad variety of cost-insensitive classifiers into cost-sensitive ones.
- MetaCost is a wrapper method which can be used for *any* type of classifier to obtain a cost-sensitive classifier. It treats the underlying classifier as a black-box in the sense that neither knowledge about its mechanism is required nor changes to its mechanism is needed.
- MetaCost needs only a cost-matrix \mathbf{C} for the underlying learning setting, which is used to adapt the decision boundaries towards predictions with low expected costs.
(Some tuning parameters are also needed.)
- Roughly speaking, the procedure of MetaCost is: relabel the training examples with their “optimal” classes, i.e., the ones with low expected costs, and apply the classifier on the relabeled data set.

METACOST: ALGORITHM

The procedure of MetaCost is divided into three phases:

- 1 **Bagging** — The underlying classifier is used (trained) L times on different bootstrapped samples of the training data, respectively.
- 2 **Relabeling** — These L trained classifiers are used to relabel the original training data set by taking the cost-matrix into account.
- 3 **Cost-sensitivity** — The classifier is trained on the relabeled data set resulting in a cost-sensitive classifier.

Predictions of the classifier f are converted (if necessary) into probabilistic prediction:

$$\text{ProbPrediction}(j, f, \mathbf{x}) = \begin{cases} (f(\mathbf{x}))_j & f \text{ is a prob. classifier,} \\ \text{One-hot}(f(\mathbf{x})) & \text{else,} \end{cases}$$

where $\text{One-hot}(f(\mathbf{x}))$ uses a one-hot-encoding of the prediction to make it a probability, i.e., one for the predicted class and 0 else.

MetaCost

Input: $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$ training data,
 $L \in \mathbb{N}$ number of bagging iterations,
 $B \in \mathbb{N}$ bootstrap size,
 f (black-box) classifier, \mathbf{C} cost matrix,
1st phase:

for $l = 1, \dots, L$ **do**

$\mathcal{D}_l \leftarrow \text{BootstrapSample}(\mathcal{D}, B)$

$f_l \leftarrow \text{train } f \text{ on } \mathcal{D}_l$

end for

2nd phase:

for $i = 1, \dots, n$ **do**

if $\mathbf{x}^{(i)} \in \mathcal{D}_l$ for all $l = 1, \dots, L$ **then**

$\tilde{L} \leftarrow \{1, \dots, L\}$

else

$\tilde{L} \leftarrow \bigcup_{l: \mathbf{x}^{(i)} \notin \mathcal{D}_l} \{l\}$

end if

for $j = 1, \dots, g$ **do** (relabel for binary case)

$p_j(\mathbf{x}^{(i)}) \leftarrow \frac{1}{|\tilde{L}|} \sum_{l \in \tilde{L}} p_j(\mathbf{x}^{(i)} | f_l)$

$p_j(\mathbf{x}^{(i)} | f_l) = \text{ProbPrediction}(j, f_l, \mathbf{x}^{(i)})$

$\tilde{y}^{(i)} \leftarrow \arg \min_{i^*} \sum_{j=1}^g p_j(\mathbf{x}^{(i)}) C(i^*, j)$

end for

$\tilde{\mathcal{D}} \leftarrow \tilde{\mathcal{D}} \cup \{(\mathbf{x}^{(i)}, \tilde{y}^{(i)})\}$

end for

3rd phase:

$f_{\text{meta}} \leftarrow \text{train } f \text{ on } \tilde{\mathcal{D}}$